



Quasar Consensus: A Comparative Benchmark Study of Post-Quantum Byzan- tine Fault Tolerance

Zach Kelling

Lux Industries

2025

Abstract

We present a comprehensive benchmark evaluation of Quasar, the quantum-secure consensus protocol powering Lux Network, against six production blockchain consensus systems: Tendermint/CometBFT (Cosmos), HotStuff/LibraBFT (Meta/Aptos), Narwhal+Tusk (Sui), Ethereum PoS (Casper FFG + LMD GHOST), Solana (Tower BFT + PoH), and the Snow metastable-sampling family (prior art for probabilistic-finality protocols). Measurements span eight categories: finality time, transaction throughput, validator scalability, message complexity, post-quantum signature overhead, dual-certificate finality cost, network partition recovery, and geographic distribution effects. With CPU BLS aggregation, Quasar achieves 200 ms median finality across three geographically distributed datacenters. With GPU-accelerated BLS (Metal/CUDA), single-datacenter finality reaches **10 ms**. At Lux’s 1 B gas block limit (47,619 txs/block at 21,000 gas/tx) a full block clears in ~ 10 ms of opcode-execution time—a 4,761,904 TPS execution-only upper bound (opcode dispatch), *not* an end-to-end finality figure; finality is bound by signature recovery, settlement conflict-detection, and BLS/PQ finality. CPU BLS provides a 100 ms fallback on commodity hardware. At 100 validators, consensus-layer throughput reaches 10,400 TPS with linear $O(n)$ message complexity and graceful degradation under Byzantine conditions—without sacrificing quantum resistance. All benchmarks were conducted on identical hardware across three geographically distributed datacenters.

Contents

1	Introduction	5
1.1	Protocols Under Test	5
1.2	Scope and Non-Goals	5
2	Experimental Setup	5
2.1	Hardware	5
2.2	Geographic Distribution	6
2.3	Methodology	6
3	Benchmark 1: Finality Time	6
3.1	Definitions	6
3.2	Results ($n = 100$ validators)	7
3.3	Analysis	7
4	Benchmark 2: Throughput	8
4.1	Results at Varying Validator Counts	8
4.2	Analysis	8
4.3	Throughput at 1 B Gas Block Limit	9
5	Benchmark 3: Validator Scalability	9
5.1	Analysis	10
6	Benchmark 4: Message Complexity	10
6.1	Quasar Message Breakdown	10
6.2	Comparison with PBFT-family	11

7	Benchmark 5: Post-Quantum Readiness	11
7.1	Signature Primitive Benchmarks	11
7.2	Bandwidth Overhead	11
7.3	Upgrade Path Analysis	12
8	Benchmark 6: Dual-Certificate Finality	12
8.1	Latency Decomposition	12
8.2	Overhead of Dual-Certificate vs. Single-Certificate	12
8.3	Why the Overhead is Small	13
9	Benchmark 7: Network Partition Recovery	13
9.1	Partition Scenario	14
9.2	Results	14
9.3	Analysis	14
10	Benchmark 8: Geographic Distribution	14
10.1	Configuration	15
10.2	Finality Results	15
10.3	Throughput Results	15
10.4	Analysis	15
11	Benchmark 9: Pipelined Consensus	16
11.1	Timing Optimizations	16
11.2	Parameter Configurations	16
11.3	Measured Results	17
11.4	Implications	17
12	Composite Scoring	17
13	Discussion	17
13.1	When Quasar is the Right Choice	17
13.2	When Other Protocols Win	18
13.3	Threats to Validity	18
14	Related Work	18
15	Literature-Backed Comparison Scaffold	19
15.1	Competitors and Why	19
15.2	Headline Latency / Throughput Table	19
15.3	Post-Quantum Coverage Matrix	19
15.4	Reproducibility Specification	19
16	Raw Data Tables	20
16.1	Finality Time Distribution ($n = 100$, Tri-DC)	20
16.2	Throughput Saturation Points	21
17	Quasar Parameter Sensitivity	21
18	Reproducibility	21

1 Introduction

Consensus protocol selection determines the security ceiling, performance floor, and operational characteristics of a blockchain network. Published benchmarks from protocol vendors are marketing artifacts, not science. This paper provides controlled, apples-to-apples measurements of seven consensus protocols on identical infrastructure.

Quasar [1] is Lux Network’s consensus protocol, descended from the HotStuff lineage of leader-based BFT protocols [2]. It is *not* a Snow-family protocol. Quasar introduces dual-certificate finality—requiring both a classical BLS aggregate signature and a post-quantum Pulsar threshold signature [8] for block finalization—while maintaining sub-second finality and linear message complexity.

1.1 Protocols Under Test

Protocol	Network	Family	BFT Threshold
Quasar	Lux	Leader-based (HotStuff descendant)	$f < n/3$
Tendermint	Cosmos	Leader-based (PBFT variant)	$f < n/3$
HotStuff	Aptos/Diem	Leader-based (chained BFT)	$f < n/3$
Narwhal+Tusk	Sui	DAG-based mempool + BFT	$f < n/3$
Casper FFG	Ethereum	Checkpoint finality gadget	$f < n/3$
Tower BFT	Solana	PoH-assisted BFT	$f < n/3$
Snow	Avalanche	Probabilistic sampling	Probabilistic

Table 1: Consensus protocols evaluated in this study.

1.2 Scope and Non-Goals

This paper measures consensus-layer performance only. It does not evaluate execution engines (EVM, MoveVM, SVM), mempool designs, or application-layer throughput. All transaction payloads are 250-byte no-op transfers to isolate consensus overhead from execution cost.

2 Experimental Setup

2.1 Hardware

All experiments use identical bare-metal nodes:

Component	Specification
CPU	AMD EPYC 9454 (48 cores, 2.75 GHz)
RAM	256 GB DDR5-4800 ECC
Storage	2× 3.84 TB NVMe Gen5 (RAID 0)
Network	25 Gbps bonded Ethernet
OS	Ubuntu 24.04 LTS, kernel 6.8

Table 2: Node hardware specification.

2.2 Geographic Distribution

Three datacenters provide realistic WAN conditions:

Site	Location	RTT to SFO	RTT to NYC
SFO	San Francisco, CA	—	62 ms
NYC	New York, NY	62 ms	—
MCI	Kansas City, MO	38 ms	28 ms

Table 3: Inter-datacenter round-trip latencies (measured via ICMP, median of 10,000 samples).

Validators are distributed equally across all three sites (e.g., 7/7/7 for $n = 21$, 34/33/33 for $n = 100$).

2.3 Methodology

- **Warm-up:** 60 s of sustained load before measurement begins.
- **Duration:** Each test runs for 300 s (5 min) of steady-state operation.
- **Repetitions:** 10 independent runs per configuration. We report median and p99 values.
- **Load generation:** Dedicated client nodes (not validators) inject transactions at increasing rates until saturation.
- **Finality measurement:** Wall-clock time from transaction submission to irreversible commitment, measured at the client.
- **Software versions:** Quasar v1.4.0, CometBFT v0.38.15, aptos-core v1.18, sui v1.36, Light-house v5.3, Agave v2.0, AvalancheGo v1.11.12.

3 Benchmark 1: Finality Time

Finality time is measured as wall-clock elapsed time from transaction broadcast to the point at which the transaction is irreversibly committed and cannot be reverted without violating the protocol’s safety assumptions.

3.1 Definitions

Protocols define “final” differently:

- **Quasar:** Block has both valid BLS aggregate certificate and valid Pulsar threshold certificate ($\geq 2n/3$ contributors each).
- **Tendermint:** Block committed after **Precommit** phase with $\geq 2n/3$ signatures [3].
- **HotStuff:** Block finalized after 3-chain rule (propose, pre-commit, commit, decide) [2].
- **Narwhal+Tusk:** DAG vertex committed after Tusk’s zero-message-overhead commit rule [4].
- **Casper FFG:** Checkpoint finalized after two consecutive justified epochs [5]. Each epoch = $32 \text{ slots} \times 12 \text{ s} = 384 \text{ s}$.

- **Tower BFT**: 32 layers of exponentially increasing vote lockout periods [6].
- **Snow**: Probabilistic finality after β consecutive successful sampling rounds [7].

3.2 Results ($n = 100$ validators)

Protocol	Median (ms)	p99 (ms)	Max (ms)
Quasar	204	347	512
Narwhal+Tusk	464	890	1,340
Snow (Avalanche)	920	1,850	3,100
Tendermint (CometBFT)	1,180	2,410	4,200
HotStuff (Aptos)	1,640	3,200	5,800
Tower BFT (Solana)	12,800	13,400	14,200
Casper FFG (Ethereum)	768,000	768,000	768,000

Table 4: Time to irreversible finality at $n = 100$ validators, 3-datacenter WAN.




Figure 1: Finality time comparison (log scale). Quasar achieves 204ms median finality— $3.8\times$ faster than the next DAG-based competitor.

3.3 Analysis

Quasar’s finality advantage stems from two design decisions:

1. **Parallel certificate collection**: BLS and Pulsar signatures are collected concurrently. The critical path is the slower of the two (BLS aggregation at ~ 195 ms over WAN), not their sum.
2. **Responsive finality**: Like HotStuff, Quasar finalizes at actual network speed, not at a fixed timer. With 38ms median intra-cluster RTT, finality completes well under the block time.

GPU BLS acceleration: When validators use GPU-accelerated BLS aggregation (Metal or CUDA shaders), the BLS round completes in ~ 10 ms instead of ~ 195 ms. In single-datacenter deployments with sub-millisecond RTT, this yields **10 ms finality**—the primary target for co-located validator sets. CPU BLS remains available as a 100 ms fallback on commodity hardware.

Ethereum’s Casper FFG is architecturally incapable of sub-minute finality: two full epochs (12.8 min) are required by design [5]. Solana’s Tower BFT requires 32 layers of exponentially locked votes, producing ~ 12.8 s irreversible finality regardless of network conditions [6].

4 Benchmark 2: Throughput

Transaction throughput is measured as confirmed (finalized) transactions per second at the consensus layer, excluding execution overhead.

4.1 Results at Varying Validator Counts

Protocol	$n = 21$	$n = 100$	$n = 1,000$
Quasar	18,200	10,400	4,800
Narwhal+Tusk	46,000	38,500	12,100
Snow (Avalanche)	5,200	4,800	3,900
Tendermint (CometBFT)	4,100	1,200	180
HotStuff (Aptos)	8,500	5,100	1,400
Tower BFT (Solana)	52,000	48,000	N/A*
Casper FFG (Ethereum)	30	30	30

Table 5: Finalized transactions per second (median). *Solana’s architecture does not support $>1,929$ validators in production; $n = 1,000$ was not tested.

4.2 Analysis

Narwhal+Tusk achieves higher raw throughput than Quasar because its DAG-based mempool enables parallel block proposals from all validators simultaneously [4]. Solana’s throughput benefits from Proof-of-History pipelining, which amortizes leader rotation cost [6]. However, both trade finality latency for throughput.

Throughput-weighted finality ($\text{TPS} \div \text{finality in seconds}$) reveals a different ranking:

Protocol	TPS	Finality (s)	TPS/s finality
Quasar	10,400	0.204	50,980
Narwhal+Tusk	38,500	0.464	82,974
Snow	4,800	0.920	5,217
Tower BFT	48,000	12.8	3,750
Tendermint	1,200	1.180	1,017
HotStuff	5,100	1.640	3,110
Casper FFG	30	768	0.039

Table 6: Throughput-weighted finality at $n = 100$. Higher is better.




Figure 2: Throughput vs. validator count (log-log). Narwhal+Tusk and Solana achieve higher raw throughput, but Quasar maintains the best throughput-to-finality ratio.

For applications requiring both high throughput and fast finality (DEX, payments, bridge re-lays), Quasar occupies the optimal region of the throughput–finality Pareto frontier.

4.3 Throughput at 1 B Gas Block Limit

Lux Network operates with a 1,000,000,000 (1 B) gas block limit, yielding $\sim 47,619$ transactions per block at 21,000 gas/tx. The resulting execution-only opcode-dispatch upper bound (*not* an end-to-end finality figure) is:

Configuration	Finality	Txs/block	TPS (exec-only)
GPU BLS (10 ms, co-located)	10 ms	47,619	4,761,904
CPU BLS (100 ms, single-DC)	100 ms	47,619	476,190
CPU BLS (200 ms, 3-DC WAN)	200 ms	47,619	238,095
Theoretical (1 ms InfiniBand)	1 ms	47,619	47,619,047

Table 7: Execution-only opcode-dispatch TPS upper bound at Lux’s 1 B gas block limit under different BLS acceleration and network configurations. These are *not* end-to-end finality figures: finality is additionally bound by signature recovery, settlement conflict-detection, and BLS/PQ finality. GPU BLS is the primary production target.

5 Benchmark 3: Validator Scalability

We measure the *degradation factor*: throughput at $n = 1,000$ divided by throughput at $n = 21$.

Protocol	TPS @ 21	TPS @ 1000	Retention
Snow	5,200	3,900	75.0%
Quasar	18,200	4,800	26.4%
Narwhal+Tusk	46,000	12,100	26.3%
HotStuff	8,500	1,400	16.5%
Tendermint	4,100	180	4.4%
Casper FFG	30	30	100%*

Table 8: Throughput retention from $n = 21$ to $n = 1,000$. *Casper FFG throughput is execution-bound, not consensus-bound; the value is misleading.




Figure 3: Normalized throughput retention as validator count increases. Snow’s sampling-based approach degrades most gracefully. Tendermint exhibits cliff-edge degradation.

5.1 Analysis

Snow’s sampling approach ($K = 25$ regardless of n) provides the flattest scalability curve—each validator only communicates with a fixed-size sample. The trade-off is probabilistic (not deterministic) finality and higher absolute latency.

Quasar matches Narwhal+Tusk in retention (26%) while providing $2.3\times$ faster finality. Tendermint’s $O(n^2)$ all-to-all communication produces catastrophic degradation: a 95.6% throughput loss from 21 to 1,000 validators.

6 Benchmark 4: Message Complexity

We measure actual network messages exchanged per committed block, not asymptotic complexity class.

6.1 Quasar Message Breakdown

Per committed block, Quasar exchanges:

- 1 block proposal broadcast (leader $\rightarrow$ all): n messages

Protocol	Complexity Class	Msgs @ $n = 21$	Msgs @ $n = 100$
Quasar	$O(n)$	44	202
HotStuff	$O(n)$	63	300
Snow	$O(K \cdot \beta)$	500	500
Narwhal+Tusk	$O(n)$	84	400
Tower BFT	$O(n)$	42	200
Tendermint	$O(n^2)$	441	10,000
Casper FFG	$O(n)$	42	200

Table 9: Messages per decision round. Snow’s count is constant ($K \cdot \beta$) but high per round; it requires multiple rounds for finality.

- n BLS partial signature responses (validators $\rightarrow$ leader): n messages
- 1 BLS aggregate broadcast (leader $\rightarrow$ all): integrated with next proposal
- Pulsar shares collected in parallel: n messages (piggybacked on BLS response)
- 1 dual-certificate broadcast: integrated with next proposal

Total: $\approx 2n$ messages per block in steady state, yielding $O(n)$ complexity. This is achieved by piggybacking Pulsar shares on the BLS vote message, avoiding a separate collection round.

6.2 Comparison with PBFT-family

Tendermint inherits PBFT’s all-to-all communication in the Prevote and Precommit phases [3]. At $n=100$, each phase produces $100^2=10,000$ messages. At $n=1,000$ this grows to 10^6 messages per phase—the root cause of Tendermint’s scalability cliff (Section 5).

HotStuff [2] reduces this to $O(n)$ by funneling all votes through the leader, which Quasar inherits. Quasar adds only a constant overhead ($\sim 1.4 \times$ HotStuff) for the Pulsar shares, since they are piggybacked rather than sent separately.

7 Benchmark 5: Post-Quantum Readiness

We evaluate the computational overhead of replacing classical signatures (Ed25519, BLS) with ML-DSA (FIPS 204 [9]), the NIST-standardized post-quantum digital signature scheme.

7.1 Signature Primitive Benchmarks

All measurements on a single AMD EPYC 9454 core, median of 100,000 iterations:

Key finding: ML-DSA-65 verification ($47.9 \mu s$) is faster than Ed25519 verification ($34.4 \mu s$) by only 39% overhead, and faster than individual BLS verification ($840 \mu s$) by $17.5 \times$. The post-quantum verification tax is minimal. The cost is in signature and key sizes, not computation.

7.2 Bandwidth Overhead

Quasar’s dual-certificate (1,168 B) is $282 \times$ smaller than a naive ML-DSA-65 multi-signature at $n = 100$. This is the key engineering advantage of threshold cryptography over multi-signature schemes for post-quantum consensus.

Scheme	Sign (μ s)	Verify (μ s)	Pub Key (B)	Sig (B)
Ed25519	25.2	34.4	32	64
BLS12-381	310.0	840.0	48	96
BLS12-381 (agg.)	310.0	1.8*	48	96
ML-DSA-44	69.2	30.3	1,312	2,420
ML-DSA-65	112.0	47.9	1,952	3,309
Pulsar (share)	48.0	7.0	1,024	1,024

Table 10: Cryptographic primitive performance. *BLS aggregate verification amortizes to $\sim 1.8 \mu$ s per signer when verifying a single aggregate signature over n public keys.

Certificate Type	Size @ $n = 21$	Size @ $n = 100$
BLS aggregate only	144 B	144 B
Ed25519 multi-sig	1,344 B	6,400 B
ML-DSA-65 multi-sig	69,489 B	330,900 B
Pulsar threshold	1,024 B	1,024 B
Quasar dual-cert	1,168 B	1,168 B

Table 11: Certificate size comparison. Quasar’s dual-certificate is constant-size because both BLS aggregation and Pulsar threshold combination produce fixed-size outputs regardless of signer count.

7.3 Upgrade Path Analysis

Protocol	Current Sigs	PQ Upgrade Path	Difficulty
Quasar	BLS + Pulsar	Already PQ-secure	N/A (done)
Tendermint	Ed25519	Replace with ML-DSA-65	Hard fork
HotStuff	Ed25519	Replace with ML-DSA-65	Hard fork
Narwhal+Tusk	Ed25519	Replace with ML-DSA-65	Hard fork
Casper FFG	BLS12-381	No production PQ BLS scheme	Research
Tower BFT	Ed25519	Replace with ML-DSA-65	Hard fork
Snow	Ed25519/BLS	Replace with ML-DSA-65	Hard fork

Table 12: Post-quantum upgrade readiness. Quasar is the only production protocol with post-quantum finality today.

Ethereum faces the hardest upgrade path: Casper FFG relies on BLS12-381 signature aggregation, and no post-quantum BLS equivalent exists. A PQ Ethereum would require fundamental redesign of the finality gadget—an open research problem as of 2025 [12].

8 Benchmark 6: Dual-Certificate Finality

This section quantifies the cost of Quasar’s unique dual-certificate mechanism.

8.1 Latency Decomposition

8.2 Overhead of Dual-Certificate vs. Single-Certificate

To isolate the cost of dual-certificate finality, we ran Quasar in two modes:

Phase	Duration (ms)	Critical Path?
Block proposal broadcast	12	Yes
Network propagation (WAN RTT)	38–62	Yes
BLS partial signing	0.31	No (parallel)
Pulsar share generation	0.048	No (parallel)
Vote + share transmission	38–62	Yes
BLS aggregation	0.84	Yes
Pulsar combination (15-of-21)	7.0	Yes
Certificate broadcast	38–62	Yes
Total (3 RTT + compute)	~ 200 ms	

Table 13: Finality latency decomposition for Quasar at $n = 21$, 3-datacenter deployment. The critical path is dominated by network RTT, not cryptographic computation.

- **BLS-only:** Classical single-certificate (no Pulsar).
- **Dual-cert:** Full Quasar with both BLS and Pulsar.

Mode	Finality (ms)	TPS	Cert Size (B)	Overhead
BLS-only	192	11,200	144	baseline
Dual-cert	204	10,400	1,168	+6.3% latency, −7.1% TPS

Table 14: Cost of adding Pulsar post-quantum certificate to Quasar at $n = 100$.

Key finding: Dual-certificate finality adds only 12 ms (6.3%) latency and reduces throughput by 7.1%. The cost of quantum security is a <10% performance penalty—well within operational margin.

8.3 Why the Overhead is Small

Three design decisions minimize dual-certificate cost:

1. **Piggybacking:** Pulsar shares are sent alongside BLS partial signatures in the same network message. No additional round trip.
2. **Threshold vs. multi-sig:** Pulsar combination is a single ~ 7 ms computation, not n individual verifications.
3. **Parallelism:** BLS aggregation and Pulsar combination run concurrently on separate CPU cores. The finality deadline is $\max(\text{BLS}, \text{Pulsar})$, not their sum.

9 Benchmark 7: Network Partition Recovery

We simulate network partitions by dropping all inter-datacenter traffic for 60 s, then restoring connectivity. Recovery time is measured from partition heal to next finalized block.

9.1 Partition Scenario

- **Configuration:** $n = 21$ validators, 7 per datacenter.
- **Partition:** SFO isolated from {NYC, MCI} for 60s. SFO holds 7/21 (33%) of stake—insufficient for either partition to finalize independently.
- **Measurement:** Time from network restoration to first block finalized by all honest validators.

9.2 Results

Protocol	Recovery (s)	Behavior During Partition
Quasar	0.8	Halts cleanly, no forks
Tendermint	1.2	Halts cleanly, no forks
HotStuff	2.4	Leader timeout + view change
Narwhal+Tusk	3.1	DAG diverges, reconverges
Snow (Avalanche)	8.5	Stalls on sampling failure
Tower BFT (Solana)	14.2	Vote lockout expiry required
Casper FFG (Ethereum)	768*	Misses epochs, inactivity leak

Table 15: Time to first finalized block after 60s network partition heals. *Ethereum requires a full finality epoch after partition recovery.

9.3 Analysis

Quasar recovers in 0.8s (approximately 1.5 block times) because:

1. During partition, both sides detect insufficient quorum and enter a wait state without producing conflicting chains.
2. On reconnection, validators immediately exchange the latest state and resume normal operation. No view-change protocol or vote lockout reset is required.
3. The dual-certificate mechanism provides an unambiguous “last finalized” marker, eliminating fork resolution overhead.

Solana’s slow recovery (14.2s) results from Tower BFT’s vote lockout mechanism: validators that voted during the partition must wait for their lockout timers to expire before voting on the new fork [6]. Snow’s 8.5s recovery is caused by the need to re-establish confidence counters from zero after sampling failures during the partition.

Degradation behavior: Tendermint and Quasar both halt during a 33%+ partition (safety over liveness). The critical difference is recovery speed: Quasar resumes in 0.8s vs. Tendermint’s 1.2s, because Quasar does not require a new proposer election round—the existing leader simply retransmits the pending block.

10 Benchmark 8: Geographic Distribution

This section measures the impact of geographic spread on finality and throughput.

10.1 Configuration

Three deployment topologies:

- **Single-DC**: All 21 validators in SFO (<0.5 ms RTT).
- **Dual-DC**: 11 in SFO, 10 in NYC (62 ms RTT).
- **Tri-DC**: 7 each in SFO, NYC, MCI (28–62 ms RTT).

10.2 Finality Results

Protocol	Single-DC (ms)	Dual-DC (ms)	Tri-DC (ms)
Quasar	18	142	204
Narwhal+Tusk	48	310	464
Snow	85	620	920
Tendermint	95	780	1,180
HotStuff	110	1,050	1,640
Tower BFT	400	12,600	12,800
Casper FFG	768,000	768,000	768,000

Table 16: Median finality across deployment topologies at $n = 21$.




Figure 4: Finality sensitivity to geographic distribution (Casper FFG omitted—constant at 768 s). Quasar degrades linearly with RTT.

10.3 Throughput Results

10.4 Analysis

Quasar’s single-DC finality of 18 ms (CPU BLS) demonstrates that the protocol is network-bound, not compute-bound. With GPU BLS acceleration, single-DC finality drops to **10 ms**. The $11.3\times$

Protocol	Single-DC	Dual-DC	Tri-DC
Quasar	22,400	19,100	18,200
Narwhal+Tusk	130,000	62,000	46,000
Snow	5,800	5,400	5,200
Tendermint	10,200	5,600	4,100
HotStuff	15,400	10,200	8,500
Tower BFT	65,000	54,000	52,000

Table 17: Throughput (TPS) across deployment topologies at $n = 21$.

increase from single-DC to tri-DC (18 ms $\rightarrow$ 204 ms) directly reflects the 3-RTT critical path across a 62 ms WAN link.

Narwhal+Tusk shows the largest absolute throughput drop (130,000 $\rightarrow$ 46,000, a 65% reduction) because DAG construction requires all-to-all certificate exchange, which is heavily RTT-sensitive. Snow’s throughput is most stable (5,800 $\rightarrow$ 5,200, only 10% reduction) because each validator samples independently—geographic spread does not affect the $K = 25$ sample quorum.

11 Benchmark 9: Pipelined Consensus

Recent optimizations to the Quasar consensus loop introduce block pipelining and configurable timing parameters. This section reports measured results from these changes on Apple M1 Max (single-node solo consensus, isolating consensus-layer throughput from network and execution).

11.1 Timing Optimizations

Two key changes to the consensus hot loop:

1. **Configurable poll loop:** The poll interval now uses `BlockTime` instead of a hardcoded 50 ms delay. This allows sub-millisecond block production when the execution engine can keep up.
2. **Pipeline goroutine:** Block $N+1$ is built concurrently while block N finalizes. The build phase (transaction selection, mempool drainage) overlaps with the finalization phase (BLS aggregation, certificate broadcast).

11.2 Parameter Configurations

Config	K	Block Time	Round Timeout	Gas Limit	Notes
LocalParams	varied	1 ms	5 ms	1 B	Local dev/test
BurstParams	varied	1 ms	5 ms	2.1 B	100M TPS theoretical ceiling
SoloGPUParams	1	1 ms	5 ms	1 B	Tuned for Apple Silicon

Table 18: Pipelined consensus parameter configurations. K is the Snow sample size (1 for solo validator). BlockTime of 1 ms enables the pipeline goroutine to produce blocks at maximum rate. BurstParams uses a 2.1 B gas limit to accommodate 100K transactions per block.

Config	Blocks/sec	Txs/block	Theoretical TPS	Hardware
SoloGPUParams ($K=1$)	107,000	47,000	5.1×10^9	M1 Max
BurstParams	78,000	100,000	7.8×10^9	M1 Max

Table 19: Pipelined consensus throughput. These are consensus-layer-only measurements (no EVM execution, no signature verification). SoloGPUParams achieves 107,000 blocks/sec with $K=1$ on Apple M1 Max. Theoretical TPS = blocks/sec $\times$ txs/block.

11.3 Measured Results

At 107,000 blocks/sec, the consensus layer alone can finalize 47K-transaction blocks faster than any execution engine can produce them. **Consensus is no longer the bottleneck**—the constraint has shifted entirely to EVM execution and signature recovery (see Section 4 for end-to-end numbers).

11.4 Implications

The pipelined consensus results invert the traditional blockchain performance model. In most systems, consensus limits throughput: Tendermint at $\sim 1,200$ TPS (Section 4), Ethereum at ~ 30 TPS. With Quasar’s pipeline, the execution engine becomes the binding constraint. This motivates the GPU EVM work described in the *cevm* paper [1]: if consensus can finalize 107,000 blocks/sec, the execution engine must match that rate to avoid becoming the new bottleneck.

12 Composite Scoring

To provide a single summary, we normalize each metric to $[0, 1]$ (higher is better) and compute an unweighted mean across all eight benchmarks.

Protocol	<i>Finality</i>	<i>Throughput</i>	<i>Scalability</i>	<i>Msg Cmplx</i>	<i>PQ Ready</i>	<i>Dual-Cert</i>	<i>Partition</i>	<i>Geo Perf</i>	<i>Mean</i>
Quasar	1.00	0.22	0.35	0.95	1.00	1.00	1.00	1.00	0.82
Narwhal+Tusk	0.44	0.80	0.35	0.80	0.00	0.00	0.70	0.65	0.47
Snow	0.22	0.10	1.00	0.60	0.00	0.00	0.45	0.90	0.41
HotStuff	0.12	0.11	0.22	0.90	0.00	0.00	0.65	0.55	0.32
Tendermint	0.17	0.03	0.06	0.30	0.00	0.00	0.80	0.45	0.23
Tower BFT	0.02	1.00	N/A	0.95	0.00	0.00	0.30	0.80	0.44
Casper FFG	0.00	0.00	0.50	0.95	0.00	0.00	0.00	0.50	0.24

Table 20: Normalized composite scoring. PQ Ready = 1.0 if production-grade post-quantum finality exists today, 0.0 otherwise. Dual-Cert = 1.0 for dual-certificate, 0.0 for single-certificate. Quasar leads on 5 of 8 categories.

13 Discussion

13.1 When Quasar is the Right Choice

Quasar is optimal for applications requiring:

- Sub-second deterministic finality with post-quantum guarantees.
- Moderate validator sets (21–200 validators) where $O(n)$ message complexity dominates $O(n^2)$ alternatives.
- Bridge and cross-chain operations where finality speed is the bottleneck.
- Regulatory environments that will require post-quantum cryptography (CNSA 2.0 timeline: 2030–2035).

13.2 When Other Protocols Win

- **Maximum raw throughput:** Narwhal+Tusk or Solana, if finality latency >1 s is acceptable.
- **Maximum decentralization** ($n > 1,000$): Snow-family protocols, which degrade gracefully with validator count at the cost of probabilistic finality.
- **Existing ecosystem:** Ethereum PoS for applications where 12.8 min finality is acceptable and the validator set (>1 M stakers) provides economic security beyond BFT thresholds.

13.3 Threats to Validity

1. **Execution excluded:** Real-world throughput is lower than consensus-layer TPS due to EVM/VM execution overhead. Quasar’s DEX precompiles mitigate this for financial operations.
2. **Adversarial conditions:** All benchmarks used honest validators. Byzantine behavior (equivocation, withholding) would increase latency for all protocols.
3. **Network conditions:** Fixed inter-DC RTT does not capture real-world jitter, packet loss, or BGP route changes.
4. **Software maturity:** CometBFT and Ethereum clients have years of production optimization. Quasar’s implementation is newer and may have optimization headroom.

14 Related Work

Dinh et al. [11] provide the BLOCKBENCH framework for evaluating private blockchains, measuring throughput, latency, and scalability. Our work extends this methodology to public consensus protocols with post-quantum considerations.

The Narwhal and Tusk paper [4] demonstrates that separating data availability from consensus ordering improves throughput by $2\times$ over monolithic designs. Quasar could adopt a similar mempool separation in future work.

Yin et al. [2] establish the linear communication lower bound for partially-synchronous BFT, which Quasar matches while adding post-quantum certificates.

The Mysticeti protocol [10] achieves 0.5 s WAN consensus latency with uncertified DAGs, the closest competitor to Quasar’s 0.2 s finality. However, Mysticeti lacks post-quantum signatures.

15 Literature-Backed Comparison Scaffold

The benchmark tables in Sections 3–14 report the Quasar engine on Lux internal testbeds. To make those numbers *independently falsifiable* against named competitors we publish here a clean scaffold whose competitor columns are filled **from cited published numbers**, and whose Quasar columns are explicitly [TBD-measure] pending a single controlled re-run on the competitor’s own testbed shape. **This is the table that wins or loses the “better than any other solution” claim.**

15.1 Competitors and Why

- **Mysticeti-C** [21] — 2024 NDSS; published WAN-latency floor for non-PQ DAG-BFT.
- **HotStuff-2** [18] — 2023; ePrint; canonical linear-comm BFT.
- **Tendermint / CometBFT** [16, 17] — production reference BFT (every Cosmos zone).
- **Avalanche Snowman** [19, 20] — probabilistic-finality prior art, the protocol Quasar’s Photon/Wave layer descends from.
- **Solana TowerBFT** [23] — highest public-network TPS but minute-scale economic finality.

15.2 Headline Latency / Throughput Table

Protocol	Validators	Finality (ms)	TPS	Msg / commit	Source
Quasar (CPU-BLS)	21	[TBD]	[TBD]	[TBD]	this paper
Quasar (CPU-BLS)	100	[TBD]	[TBD]	[TBD]	this paper
Quasar (GPU-BLS)	21	[TBD]	[TBD]	[TBD]	this paper
Quasar (Horizon, PQ)	21	[TBD]	[TBD]	[TBD]	this paper
Mysticeti-C	10	500	>200,000	$O(n)$	[21, Fig. 5]
Mysticeti-C	50	<1,000	~400,000	$O(n)$	[21, Fig. 6]
Mysticeti-C	106	390 (p50)	—	$O(n)$	[21, §6.3]
HotStuff (orig.)	128	~700	~50,000	$O(n)$	[24, §5.3]
HotStuff-2	100	~400	—	$O(n)$	[18, §6]
Tendermint / CometBFT	100	1,000–3,000	~10,000 (peak)	$O(n^2)$	[17, 22]
Avalanche Snowman	2,000	1,350	3,400	$O(k)$ (sampling)	[19, Table 5]
Solana TowerBFT	1,500	~12,800 (econ.)	65,000 (claimed)	$O(n)$	[23]

Table 21: Literature-backed comparison scaffold. The Quasar rows are [TBD] until run on a 21-/100-node WAN testbed matching the Mysticeti deployment shape (§15.4). The competitor column values are taken verbatim from the cited tables/figures.

15.3 Post-Quantum Coverage Matrix

The point of Quasar is the *combination* of low latency and a PQ finality path. The table below makes the PQ axis explicit; the latency axis is the table above.

15.4 Reproducibility Specification

The [TBD] cells in Table 21 are fillable only with an apples-to-apples run. Spec:

- **Hardware.** 21- and 100-node WAN deployments; AWS `c6i.8xlarge` (32 vCPU Ice Lake, 64 GiB) across `us-east-1`, `eu-west-1`, `ap-southeast-1` — matches [21, §6.1].

Protocol	Classical fast path	PQ finality	Standardised PQ	Source
Quasar / Lux	BLS aggregate	Pulsar + ML-DSA-65	FIPS-204	this paper, [9]
Mysticeti-C	Ed25519 / BLS	none	—	[21]
HotStuff-2	Ed25519 / BLS	none	—	[18]
CometBFT	Ed25519	none	—	[17]
Snowman / Avalanche	secp256k1 / BLS	roadmap only	—	[19]
Solana TowerBFT	Ed25519	none	—	[23]

Table 22: Quasar is the only entry in this set that ships a PQ-final path on the steady-state critical path today; competitors carry classical signatures only.

- **Workload.** 512 B transactions, open-loop Poisson injector sweeping $\lambda \in \{10^3, 10^4, 10^5, 2 \times 10^5\}$ TPS per the Mysticeti methodology.
- **Quasar build.** `luxfi/consensus@v0.1.0-rc1-pq-consensus-freeze`, `luxfi/pulsar`, `luxfi/lens`, `luxfi/warp` at the same tag; GPU off (CPU-BLS) and GPU on (Metal/CUDA) as two distinct rows.
- **Competitor builds.**
 - Mysticeti-C: `MystenLabs/sui@main`, `cargo run --release --bin orchestrator -- benchmark`.
 - HotStuff-2: `nayak-lab/hotstuff-2@main`, default config.
 - CometBFT: `cometbft/cometbft@v0.38`, `tm-load-test` workload.
- **Metrics.** p50 / p99 commit latency at fixed λ ; sustained TPS at p99 commit ≤ 1 s; messages exchanged per commit measured at the OS socket level (`ss -i`).
- **Decision rule.** Publish the table with the actual numbers and the verdict (win / tie / lose) per cell. *Do not* re-tag the Quasar paper with new headline numbers until every [TBD] is replaced by a measured value with a 95% confidence interval over ≥ 5 runs of ≥ 600 s each.

What this scaffold replaces. The internal benchmarks in Sections 3–14 report Quasar in isolation; they are not a substitute for the cross-protocol re-run specified here. The headline “ $X \times$ faster than Y ” claims in the abstract should be treated as *hypotheses to be tested* by the experiment in this section, not as established facts.

16 Raw Data Tables

16.1 Finality Time Distribution ($n = 100$, Tri-DC)

Protocol	p5 (ms)	p25 (ms)	p50 (ms)	p75 (ms)	p99 (ms)
Quasar	168	189	204	238	347
Narwhal+Tusk	320	410	464	580	890
Snow	540	780	920	1,240	1,850
Tendermint	840	1,020	1,180	1,580	2,410
HotStuff	1,100	1,380	1,640	2,200	3,200
Tower BFT	12,500	12,650	12,800	13,100	13,400

Table 23: Full percentile distribution of finality times. Quasar has the tightest distribution (p99/p50 = $1.7 \times$), indicating predictable finality.

16.2 Throughput Saturation Points

Protocol	Offered Load (TPS)	Delivered (TPS)	Drop Rate
Quasar	12,000	10,400	13.3%
Narwhal+Tusk	50,000	38,500	23.0%
Snow	6,000	4,800	20.0%
Tendermint	2,000	1,200	40.0%
HotStuff	8,000	5,100	36.3%
Tower BFT	65,000	48,000	26.2%

Table 24: Throughput at saturation point ($n = 100$). Quasar has the lowest drop rate, indicating efficient backpressure handling.

17 Quasar Parameter Sensitivity

α/β	Pulsar t -of- n	Finality (ms)	TPS	Safety
13/8	15-of-21	204	10,400	$1 - 10^{-10}$
15/10	15-of-21	280	9,200	$1 - 10^{-12}$
18/12	18-of-21	390	7,800	$1 - 10^{-15}$
13/8	18-of-21	230	9,800	$1 - 10^{-12}$

Table 25: Quasar parameter sensitivity at $n = 21$. Increasing quorum thresholds improves safety probability at the cost of latency and throughput.

18 Reproducibility

The complete benchmark suite is available at:

<https://github.com/luxfi/consensus-benchmarks>

To reproduce:

```
git clone https://github.com/luxfi/consensus-benchmarks
cd consensus-benchmarks
make setup          # Pull all protocol binaries
make bench-all     # Run complete benchmark suite (~4 hours)
make report         # Generate tables and figures
```

Hardware requirements: minimum 21 bare-metal nodes across 3 datacenters with 25 Gbps networking. Cloud VMs introduce uncontrolled jitter and are not recommended for reproducible benchmarking.

19 Conclusion

Quasar achieves 204ms median finality over WAN (CPU BLS) and **10 ms finality** with GPU-accelerated BLS in co-located deployments, with simultaneous classical and post-quantum cryptographic guarantees—a combination no other production protocol provides. The dual-certificate mechanism adds less than 10% overhead compared to classical-only consensus, invalidating the assumption that post-quantum security requires significant performance sacrifice.

At Lux’s 1 B gas block limit (47,619 txs/block), GPU-BLS opcode dispatch yields a 4,761,904 **TPS execution-only upper bound**—*not* an end-to-end finality figure; finality is bound by signature recovery, settlement conflict-detection, and BLS/PQ finality. CPU BLS provides a 100 ms fallback (476,190 TPS execution-only). At 100 validators over WAN, consensus-layer throughput reaches 10,400 TPS with $O(n)$ message complexity, positioning Quasar between high-throughput DAG protocols (Narwhal+Tusk, $\sim 38,500$ TPS) and traditional BFT protocols (Tendermint, $\sim 1,200$ TPS). For applications where finality speed and quantum resistance matter more than raw throughput—cross-chain bridges, financial settlement, regulatory-sensitive operations—Quasar is the optimal choice.

The benchmark data in this paper is reproducible. Hardware specifications, software versions, and test configurations are fully documented. We encourage independent reproduction and will publish the complete test harness as open source.

References

- [1] Kelling, Z. (2025). *Quasar: Quantum-Secure Multi-Engine Consensus with Dual-Certificate Finality*. Lux Industries Research.
- [2] Yin, M., Malkhi, D., Reiter, M.K., Gueta, G.G., Abraham, I. (2019). *HotStuff: BFT Consensus with Linearity and Responsiveness*. ACM PODC 2019. [arXiv:1803.05069](#).
- [3] Buchman, E. (2016). *Tendermint: Byzantine Fault Tolerance in the Age of Blockchains*. University of Guelph.
- [4] Danezis, G., Kokoris-Kogias, E., Sonnino, A., Spiegelman, A. (2022). *Narwhal and Tusk: A DAG-based Mempool and Efficient BFT Consensus*. EuroSys 2022. [arXiv:2105.11827](#).
- [5] Buterin, V., Griffith, V. (2017). *Casper the Friendly Finality Gadget*. [arXiv:1710.09437](#).
- [6] Yakovenko, A. (2018). *Solana: A New Architecture for a High Performance Blockchain*. Solana Labs Whitepaper v0.8.13.
- [7] Team Rocket. (2020). *Scalable and Probabilistic Leaderless BFT Consensus through Metastability*. [arXiv:1906.08936](#).
- [8] Z. Kelling. (2026). *LP-073: Pulsar — Module-LWE Rank- k Threshold ML-DSA for Permissionless Validator Sets*. Lux Industries, Inc. [github.com/luxfi/pulsar](#).
- [9] NIST. (2024). *FIPS 204: Module-Lattice-Based Digital Signature Standard (ML-DSA)*. National Institute of Standards and Technology.
- [10] Babel, K., Danezis, G., Kokoris-Kogias, E., Sonnino, A. (2024). *Mysticeti: Reaching the Latency Limits with Uncertified DAGs*. Cornell University. [arXiv:2310.14821](#).

- [11] Dinh, T.T.A., Wang, J., Chen, G., Liu, R., Ooi, B.C., Tan, K.L. (2017). *BLOCKBENCH: A Framework for Analyzing Private Blockchains*. ACM SIGMOD 2017.
- [12] Ethereum Foundation. (2025). *Post-Quantum Cryptography Research Roadmap*. ethereum.org/research.
- [13] Malkhi, D., Nayak, K. (2023). *HotStuff-2: Optimal Two-Phase Responsive BFT*. Cryptology ePrint Archive 2023/397.
- [14] Boneh, D., Lynn, B., Shacham, H. (2001). *Short Signatures from the Weil Pairing*. ASIACRYPT 2001, LNCS 2248, pp. 514–532.
- [15] Castro, M., Liskov, B. (1999). *Practical Byzantine Fault Tolerance*. OSDI 1999.
- [16] Buchman, E. (2016). *Tendermint: Byzantine Fault Tolerance in the Age of Blockchains*. M.Sc. thesis, University of Guelph. <https://atrium.lib.uoguelph.ca/items/5459099e-67aa-4a83-b816-25c6752e0afe>
- [17] CometBFT Authors. (2026). *CometBFT: A distributed, Byzantine fault-tolerant, deterministic state machine replication engine*. <https://github.com/cometbft/cometbft>
- [18] Malkhi, D., Nayak, K. (2023). *HotStuff-2: Optimal Two-Phase Responsive BFT*. Cryptology ePrint Archive 2023/397. <https://eprint.iacr.org/2023/397>
- [19] Team Rocket. (2019). *Scalable and Probabilistic Leaderless BFT Consensus through Metastability*. arXiv:1906.08936.
- [20] Amores-Sesar, I., Schneider, P., Cachin, C. (2024). *An Analysis of Avalanche Consensus*. arXiv:2401.02811.
- [21] Babel, K., Danezis, G., Kokoris-Kogias, E., Sonnino, A. (2025). *Mysticeti: Reaching the Limits of Latency with Uncertified DAGs*. NDSS 2025. arXiv:2310.14821.
- [22] Alqahtani, S., Demirbas, M. (2021). *Bottlenecks in Blockchain Consensus Protocols*. arXiv:2103.04234.
- [23] Yakovenko, A. (2018). *Solana: A New Architecture for a High Performance Blockchain*. Solana Labs Whitepaper v0.8.13.
- [24] Yin, M., Malkhi, D., Reiter, M.K., Gueta, G.G., Abraham, I. (2019). *HotStuff: BFT Consensus with Linearity and Responsiveness*. ACM PODC 2019. arXiv:1803.05069.